


Redundancy-aware visual-token selection used by Ours

The SparseVLM relevance signal is preserved, then selected tokens are diversified before inference continues.



Applied at SparseVLM pruning layers while keeping the original model and token-recycling pipeline unchanged.

α : candidate-pool factor; λ : relevance weight; k_ℓ : layer-wise sparse token budget.