


Per-sample inference latency by method

GQA



POPE

